

In [2]: 1 !nvidia-smi

Mon Feb 1 16:50:04 2021									
+-----+									
NVIDIA-SMI		460.32.03		Driver Version:		460.32.03		CUDA Version: 11.2	
+-----+									
GPU Name		Persistence-M		Bus-Id		Disp.A		Volatile Uncorr. ECC	
Fan Temp Perf		Pwr:Usage/Cap		Memory-Usage		GPU-Util		Compute M.	
								MIG M.	
+-----+									
0	GeForce GTX 108...	Off		00000000:18:00.0	Off			N/A	
70%	80C	P2	99W / 250W	10809MiB / 11178MiB		100%	Default	N/A	
+-----+									
1	GeForce GTX 108...	Off		00000000:3B:00.0	Off			N/A	
0%	47C	P8	10W / 250W	2453MiB / 11178MiB		0%	Default	N/A	
+-----+									
2	GeForce GTX 108...	Off		00000000:5E:00.0	Off			N/A	
0%	39C	P8	9W / 250W	5887MiB / 11178MiB		0%	Default	N/A	
+-----+									
3	GeForce GTX 108...	Off		00000000:86:00.0	Off			N/A	
68%	80C	P2	249W / 250W	10809MiB / 11178MiB		78%	Default	N/A	
+-----+									
+-----+									
Processes:									
GPU	GI	CI	PID	Type	Process name	GPU Memory			
	ID	ID				Usage			
+-----+									
0	N/A	N/A	1730	G	/usr/lib/xorg/Xorg	4MiB			
0	N/A	N/A	905721	C	python	10801MiB			
1	N/A	N/A	1730	G	/usr/lib/xorg/Xorg	4MiB			
1	N/A	N/A	577506	C	...khov/anaconda3/bin/python	991MiB			
1	N/A	N/A	1449170	C	...a3/envs/common/bin/python	713MiB			
1	N/A	N/A	2155081	C	...khov/anaconda3/bin/python	741MiB			
2	N/A	N/A	1730	G	/usr/lib/xorg/Xorg	4MiB			
2	N/A	N/A	364307	C	...a3/envs/common/bin/python	711MiB			
2	N/A	N/A	956198	C	...a3/envs/common/bin/python	669MiB			
2	N/A	N/A	1751671	C	...a3/envs/common/bin/python	1471MiB			
2	N/A	N/A	2813905	C	...a3/envs/common/bin/python	2227MiB			
2	N/A	N/A	3154363	C	...a3/envs/common/bin/python	801MiB			
3	N/A	N/A	1730	G	/usr/lib/xorg/Xorg	4MiB			
3	N/A	N/A	1982195	C	python	10801MiB			
+-----+									

In [3]: 1 # from models import Encoder, Decoder, Img2Caption
2 from dataset import CocoDataset

In [3]: 1 import torchvision
2 import torch
3 import torch.nn as nn
4 import torch.nn.functional as F
5 from torch.utils.data import DataLoader
6
7 import numpy as np
8
9 from pycocotools.coco import COCO
10 from nltk.translate.bleu_score import sentence_bleu
11
12 import tqdm
13 import matplotlib.pyplot as plt
14 from IPython.display import clear_output
15
16 from transformers import DebertaTokenizer, DebertaModel

In [4]: 1 import os
2 os.environ["CUDA_VISIBLE_DEVICES"]="1"
3 torch.cuda.empty_cache()

In [6]: 1 # training images
2 len(os.listdir("images_train/"))

Out[6]: 14117

In [7]: 1 # val images
2 len(os.listdir("images_val/"))

Out[7]: 8376

In [6]: 1 # current dir
2 os.listdir()

Out[6]: ['.idea',
'ipynb_checkpoints',
'annotations',
'best_val_model.pt',
'dataset.py',
'dump_1.pt',
'images_train',
'images_val',
'kot.jpg',
'models.py',
'Training.ipynb',
'Usage example.ipynb',
'__pycache__']

In [8]: 1 # coco = COCO(annotation_file="annotations/captions_train2014.json")
2 # coco.download("images_train", imgIds=[coco.anns[key]['image_id'] for key in coco.anns])

In [9]: 1 # coco = COCO(annotation_file="annotations/captions_val2014.json")
2 # coco.download("images_val", imgIds=[coco.anns[key]['image_id'] for key in coco.anns])

In [10]: 1 tokenizer = DebertaTokenizer.from_pretrained("microsoft/deberta-base")

In [11]: 1 device = "cuda"

```
In [12]: 1 def collate_fn(data):
2         images, captions = zip(*data)
3
4         images = torch.stack(images, 0)
5         max_len = max([len(caption) for caption in captions])
6         targets = torch.zeros(len(captions), max_len).long()
7         for i, caption in enumerate(captions):
8             for j in range(len(caption)):
9                 targets[i, j] = caption[j]
10
11         return images.to(device), targets.to(device)
```

```
In [13]: 1 batch_size = 16
2 train_data = CocoDataset("images_train/", "annotations/captions_train2014.json", tokenizer)
3 test_data = CocoDataset("images_val/", "annotations/captions_val2014.json", tokenizer)
4
5 train_loader = DataLoader(train_data,
6                           batch_size=batch_size,
7                           collate_fn=collate_fn)
8 test_loader = DataLoader(test_data,
9                           batch_size=batch_size,
10                          collate_fn=collate_fn)
```

```
loading annotations into memory...
Done (t=0.82s)
creating index...
index created!
loading annotations into memory...
Done (t=0.37s)
creating index...
index created!
```

```
In [14]: 1 class Decoder(nn.Module):
2         def __init__(self, vocab_size, embedding_dim, encoder_dim, attention_dim, hidden_size, dropout):
3             super().__init__()
4             self.encoder_dim = encoder_dim
5             self.embedding_dim = embedding_dim
6             self.hidden_size = hidden_size
7
8             self.attention_dim = attention_dim
9             self.vocab_size = vocab_size
10
11            self.attention = nn.Linear(encoder_dim+embedding_dim, attention_dim)
12
13            self.rnn = nn.GRU(attention_dim, hidden_size)
14
15            self.out = nn.Linear(hidden_size+attention_dim, vocab_size)
16
17            self.dropout = nn.Dropout(dropout)
18
19            def forward(self, inputs, hidden, encoder_outputs):
20                concated_inputs = torch.cat((inputs, encoder_outputs.unsqueeze(1)), dim=2)
21                attended = self.dropout(torch.tanh(self.attention(concated_inputs))).permute(1, 0, 2)
22
23                output, hidden = self.rnn(attended, hidden)
24                concated_outputs = F.relu(torch.cat((output, attended), dim=2))
25                out = self.out(self.dropout(concated_outputs))
26
27                return out, hidden
28
29            def initHidden(self, bs):
30                return torch.zeros(1, bs, self.hidden_size, device=device)
```

```
In [15]: 1 class Img2Caption(nn.Module):
2         def __init__(self, encoder, decoder, gptModel):
3             super().__init__()
4
5             self.encoder = encoder
6             self.decoder = decoder
7             self.gpt = gptModel
8
9            def forward(self, image, caption, teacher_forcing_ratio=0.5):
10            # caption: bs x max_len
11            batch_size = caption.shape[0]
12            max_len = caption.shape[1]
13            trg_vocab_size = self.decoder.vocab_size
14            outputs = torch.zeros(max_len, batch_size, trg_vocab_size)
15            encoder_outputs = self.encoder(image)
16            caption = gptModel(caption).last_hidden_state
17
18            first_input = caption[:, 0].unsqueeze(1)
19            hidden = decoder.initHidden(first_input.shape[0])
20
21            for t in range(1, max_len):
22                output, hidden = self.decoder(first_input, hidden, encoder_outputs)
23                outputs[t] = output
24
25                teacher_force = np.random.random() < teacher_forcing_ratio
26                if teacher_force:
27                    first_input = caption[:, t].unsqueeze(1)
28                else:
29                    first_input = self.gpt(torch.argmax(output, dim=2)).last_hidden_state.permute(1, 0, 2)
30
31            return outputs
```

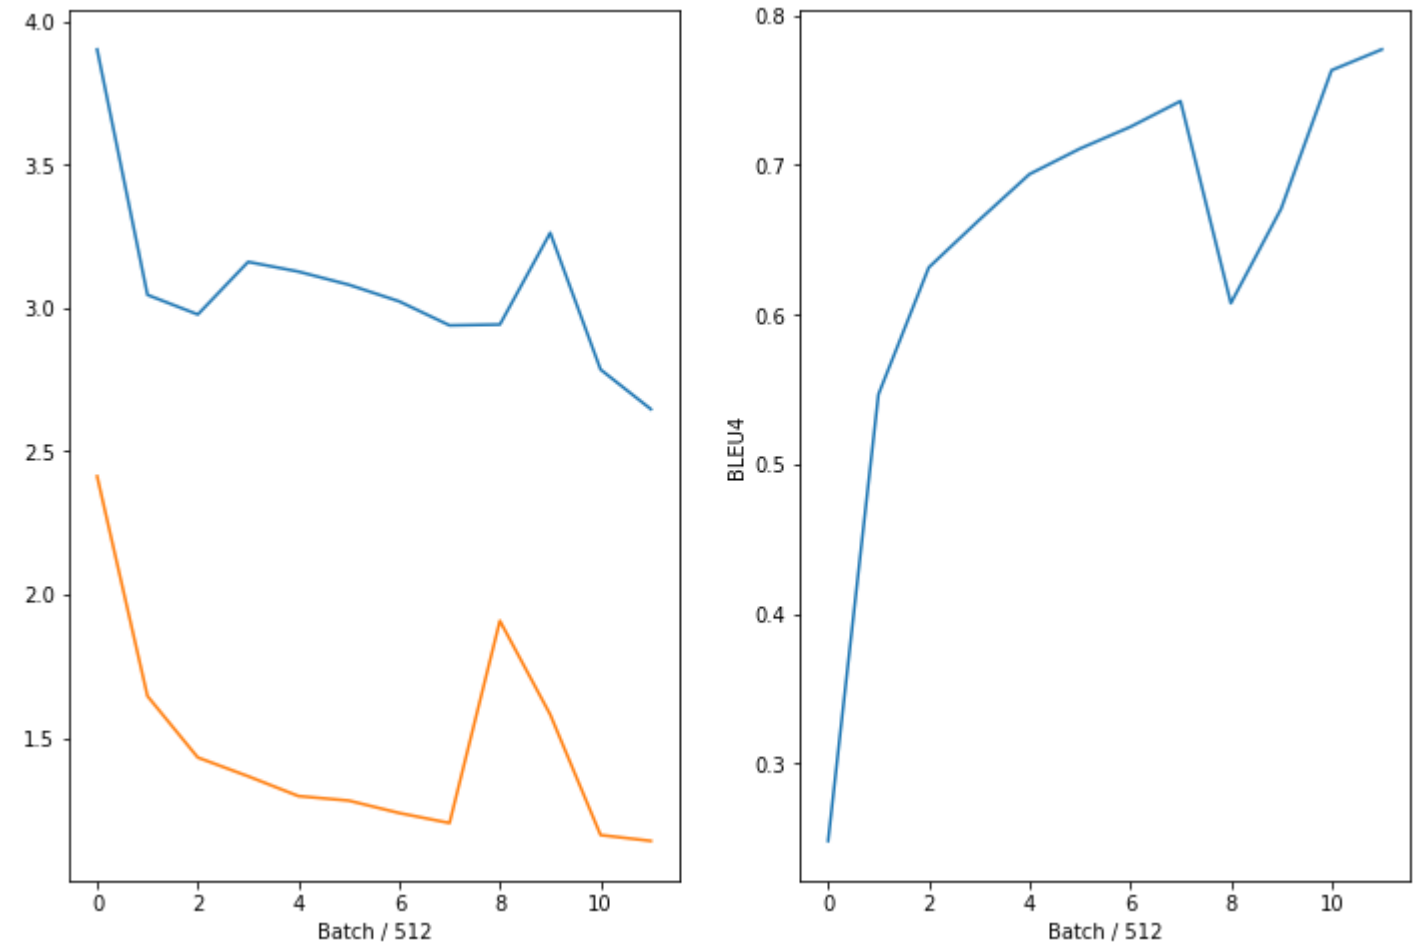
```
In [16]: 1 encoder_dim = 512
2 model_conv=torchvision.models.resnet101(pretrained=True)
3 model_conv.fc = nn.Linear(in_features=2048, out_features=encoder_dim)
4
5 for i, pair in enumerate(model_conv.named_children()):
6     _, child = pair
7     if len(list((model_conv.named_children()))) - i > 3:
8         for _, params in child.named_parameters():
9             params.requires_grad = False
10 model_conv = model_conv.to(device)
```

```
In [17]: 1 # vocab_size, embedding_dim, encoder_dim, attention_dim, hidden_size, decoder_dim, dropout
2 decoder = Decoder(tokenizer.vocab_size, 768, encoder_dim, 256, 256, 0.2).to(device)
3 gptModel = DebertaModel.from_pretrained('microsoft/deberta-base').to(device)
4
5 model = Img2Caption(model_conv, decoder, gptModel).to(device)
6
7 criterion = nn.CrossEntropyLoss()
8 optimizer = torch.optim.Adam(model.parameters(), lr=0.0001)
```

Some weights of the model checkpoint at microsoft/deberta-base were not used when initializing DebertaModel: ['deberta.embeddings.position_embeddings.weight']
- This IS expected if you are initializing DebertaModel from the checkpoint of a model trained on another task or with another architecture (e.g. initializing a BertForSequenceClassification model from a BertForPreTraining model).
- This IS NOT expected if you are initializing DebertaModel from the checkpoint of a model that you expect to be exactly identical (initializing a BertForSequenceClassification model from a BertForSequenceClassification model).

```
In [18]: 1 def decode_captions(captions):
2         value = tokenizer.decode(captions)
3         print(value)
4         return value
5
6 def calculate_bleu(captions, targets):
7     captions = torch.argmax(captions.permute(1, 0, 2), axis=2)
8     cum_bleu = 0
9     for i in range(len(captions)):
10        cum_bleu += sentence_bleu(
11            [decode_captions(targets[i]).split()],
12            decode_captions(captions[i]).split()
13        )
14    print(cum_bleu)
15    return cum_bleu / len(captions)
```

```
In [ ]: 1 %matplotlib inline
2 model.train()
3 torch.cuda.empty_cache()
4
5 logging_freq = 512
6
7
8 max_loss = np.inf
9
10 train_loss_accum = 0
11 history = []
12 val_bleu_history = []
13 val_loss_history = []
14 for epoch in range(20):
15     for i, batch in enumerate(tqdm.notebook.tqdm(train_loader)):
16
17         images, captions = batch
18
19         optimizer.zero_grad()
20         output = model(images, captions, teacher_forcing_ratio=1/np.log(2+epoch))
21         loss = criterion(output.permute(1, 0, 2).reshape(-1, tokenizer.vocab_size).cpu(), captions.view(-1).cpu())
22         loss.backward()
23         torch.nn.utils.clip_grad_norm_(model.parameters(), 1)
24
25         optimizer.step()
26         train_loss_accum += loss
27         if (i+1)%logging_freq==0:
28             history.append(train_loss_accum / logging_freq)
29             train_loss_accum = 0
30
31         loss_accum = 0
32         bleu_accum = 0
33         for test_batch in tqdm.notebook.tqdm(test_loader):
34             images, captions = test_batch
35             output = model(images, captions, teacher_forcing_ratio=1)
36             test_loss = criterion(output.permute(1, 0, 2).reshape(-1, tokenizer.vocab_size).cpu(), captions.view(-1).cpu())
37             loss_accum += test_loss.cpu().data.numpy()
38             bleu_accum += calculate_bleu(output, captions)
39         val_bleu_history.append(bleu_accum / len(test_loader))
40         val_loss_history.append(loss_accum / len(test_loader))
41
42         clear_output(True)
43         fig, ax = plt.subplots(nrows=1, ncols=2, figsize=(12, 8))
44
45
46         ax[0].plot(history, label='train loss')
47         if val_loss_history is not None:
48             ax[0].plot(val_loss_history, label='validation loss')
49             ax[0].set_xlabel(f'Batch / {logging_freq}')
50         if val_bleu_history is not None:
51             ax[1].plot(val_bleu_history, label='validation bleu history')
52             ax[1].set_xlabel(f'Batch / {logging_freq}')
53             ax[1].set_ylabel('BLEU4')
54         # plt.legend()
55         plt.show()
56
57
58         if val_loss_history[-1] < max_loss:
59             max_loss = val_loss_history[-1]
60             torch.save(model.state_dict(), 'best_val_model.pt')
```



```
HBox(children=(HTML(value=''), FloatProgress(value=0.0, max=883.0), HTML(value='')))
```

In [20]:

1 # bleu on train_split

2 calculate_bleu(output, captions)

A black stove sitting between counters in a kitchen.
A kitchen stove with a counters. a.
A clean kitchen with a bar and large window.
A kitchen kitchen with a bar. a..
A tile bathroom with a toilet, cabinet, sink, mirror, tub, shower curtain, and towel is shown.
A bathroom bathroom with a toilet a a and sink. mirror,. and is.
a white kitchen a table and a white refrigerator
a kitchen kitchen a a and a a.
A man is chopping green peppers with a knife.
A man is of a chairs. a.
A bike is parked on a bridge without a person.
A man is a a a a.. person.
A large kitchen with a tiled floor and two refrigerators.
A kitchen kitchen with a t a a a two. kitchen.
A dog is standing and looking through the threshold into a room with sinks and and a mirror
A bathroom is a a looking in a a into a room. and
A man stands in the bathroom in front of the sink.
A kitchen stands with a bathroom with a the.
A bicycle lying on a beach with the sunset in the background.
A man lying on a beach on a in a background.
An arrangement of utensils sitting outside on a table.
An kitchen of a withils in a a a.
A newly modeled kitchen with a bicycle poster.
A kitchen bathroom with a a kitchen. a
Several people are spending time in a park with a jet airliner on the ground in the background
Several a are a a in a a on a. a. ground.
a police man in uniform is riding a red motorcycle
a man man a a is in a a motorcycle.
a small plane on an air port run way
a man plane on a air. a
A bathroom with a large bathtub and a sink.
A bathroom with a bathroom bath with a a sink.
A person getting ready to board a bus with a man on a bike in front of it.
A around a a board a a. a on a in a.
A kitchen with lots of counter space next to a window.
A kitchen with a with counter a a a.
A countertop holding a chicken and a bundt cake.
A kitchentop on a a. a..
some girls are in a kitchen and one is doing dishes
some man are with in kitchen in a doing.
a woman in a long coat and her pink umbrella
a man in a a flowers on a umbrella
A bus has bicycles mounted on the front.
A man has a a on a a.
A kitchen sink next to a window and an automatic dishwasher.
A kitchen sink with a a with. a a.washer
There are several people enjoying the fresh air and sunshine in this field.
There of several a a the with a a a. this..
A granite counter top with a sink is centered in a kitchen with a stove, cabinets, microwave, and a framed vintage advertisement for bicycles on the wall in the background.
A kitchen counter kitchen kitchen a kitchen. a in a kitchen. stove. microwave and a field. for. on the wall in the kitchen.
A bathroom with a mirror and decorations on the wall.
A bathroom with a kitchen and a and a wall.
A bathroom sink beneath a very large mirror reflecting a roll of toilet paper.
A bathroom sink and a very a a a a bathroom of a a.
A large group of bicycles are in a row.
A bathroom group a a are a a..
Three people sitting on a wooden bench eating plates of food.
Three people sitting in a wooden a a. of a.
A man is asleep with a colt on a mattress.
A man is on a a a in a a.
Three men standing around a kitchen having a drink together.
Three man standing in a kitchen a a a together.
A large bathroom with tile flooring and white fixtures
A bathroom bathroom with a floor. a
1.0122869976420739

Out[20]: 0.03163396867631481

Отчет

О том что есть в этом ноутбуке:

- 1. Что-то явно случилось с валидационной выборкой - почему-то на ней запредельный bleu4, а на трейне bleu4 в несколько раз меньше. Кажется, что в дальнейшем стоит удалить данные и скачать их заново.
- 2. Модель недообучена: я долго дебажил вывод модельки, а resnet101 + deberta + декодер занимают 8гб видеопамяти - спустя неделю меня выгнали с кластера.
- 3. Файлы датасета и модели отдельно вынесены в dataset.py/models.py (здесь оставлены для удобства проверки).

О том чего нет в ноутбуке:

- 1. Изначально пробовал бре токенизацию из youtokentome, не особо получилось
- 2. Пробовал w2v/обычный берт эмбединги - была бага в препроцессинге текста, поэтому с ними результат не зафиксировал
- 3. Пробовал прикрутить "умный" attention с разных слоев resnet - не очень получилось
- 4. Пример капшининга одной функцией в ноутбуке Usage example

In []:

1